



Curtin University

5/1/2016

OS200 Assignment Report

BEARDSMORE, CONNOR - 15504319

Practical time: THURSDAY 5:00pm
OPERATING SYSTEMS 200

Source Code

The software solution for this assignment - including all source code - is included in Appendix A. The source code is written in the C language and follows the C99 standard. Additionally, attached in Appendix B is the readme file associated with the source code, explaining how to correctly run and compile the program.

Mutual Exclusion

Mutual exclusion was achieved in different ways for the multiprocessing and the multithreaded solutions. The multiprocessing solution utilized three POSIX semaphores (`sem_t`). The first represented a standard mutex lock, while the others represented full and empty flags. These semaphores ensured no race conditions occurred and that all access to shared memory was synchronized efficiently and properly. The following code segment indicates the basic locking sequence utilized by both the producer and consumer processes:

Producer:

```
wait(empty);
wait(mutex);
    // critical section
signal(mutex);
signal(full);
```

Consumer:

```
wait(full);
wait(mutex);
    // critical section
signal(mutex);
signal(empty);
```

The multithreaded solution also employed the use of three locks. The first was a POSIX mutex lock, while the others were POSIX conditions representing the "full" and "empty" conditions. These conditions deny "hold and wait" by forcing a thread to give up its mutex lock while waiting for the condition. Without "hold-and-wait", deadlock is not possible in this situation. The following code segment indicates the basic locking sequence utilized by both the producer and consumer threads:

Producer:

```
lock(mutex)
cond_wait(full)    // implicitly unlock mutex
    // critical section
cond_signal(empty)
unlock(mutex)
```

Consumer:

```
lock(mutex)
cond_wait(empty)    // implicitly unlock mutex
    // critical section
cond_signal(full)
unlock(mutex)
```

Shared Memory

Threads share the address space of the process they are created within and thus require no shared memory to be established. For the threading solution, variables such as the matrices, their dimensions and the subtotal struct were simply declared globally, so each thread could access this data without additional function import overheads. In the multiprocessing solution, POSIX shared memory was employed via the four functions below:

```
shm_open()  
ftruncate()  
mmap()  
close()
```

Unlike threads, child processes do not inherit their parents address space on creation. As a result, shared memory is required for both the parent and the child to access the same data concurrently. All three matrices were created in a separate shared memory block, as was the subtotal struct and the struct containing the semaphore variables. The child processes only ever read from the first and second matrices and thus there was no issues with concurrent access to this data. The product matrix differed from the other shared memory segments in that only one process ever writes to one specific row. Hence, no synchronization was required to read or write to any of the matrices.

Zombie Processes

The creation of zombie processes was a potential issue in the multiprocessing solution. The parent was required to perform work as the consumer and thus, could not call wait on its child to reap it normally upon termination. To avoid the prevention of zombie processes, the signal() function is called prior to forking with both the SIGCHLD and SIG_IGN flags. Any child process that terminates after explicitly setting these flags will be immediately removed from the system and thus, no zombie processes are ever formed.

The downside of this solution is that the exit status of the child process is ignored. Due to the scale of this program, it was not an issue. An alternative solution could have been to implement double forking. However, this dramatically slows down the program as double the child processes are essentially forked.

Errors and Bugs

I am currently unaware as to any error conditions that would cause either the multithreaded or multiprocessing solution to fail. There are limits however to the maximum number of processes and threads that can physically be handled. The maximum limit of processes that can be handled on the tested system was approximately 800 and the maximum limit of threads was approximately 300. This limitation is likely due to the available system resources of the laboratory computers. There have also been instances of the *shm_open()* function failing to create shared memory segments on *ssh* access to the lab computers. This however is a fault of the servers themselves and not of the program itself.

While there is primitive error checking throughout my program, more technical error conditions such as certain file input formats can result in unspecified behaviour. All shared memory functionality and use of POSIX locking variables has been error checked to a relevant level for the scale of this program.

Testing

Both solutions have been thoroughly tested on the laboratory computers in 314.219 and on *ssh* access to these computers via "*saeshell*". Functions were developed to print the entire contents of the matrices to enable debugging and to ensure that all files were read properly into the matrices. The subsequent table lists the types of all test files the programs were tested on prior to submission:

File Name	Dimensions	Contents	Expected Total
allOnes.txt	100 x 100	Every element = 1	1,000,000
exampleA exampleB	3 x 2 2 x 4	Example from specifications	402
medium A medium B	10 x 10 10 x 10	Randomly generated numbers, 1 - 10	27,132
largeA largeB	30 x 10 10 x 30	Randomly generated numbers, 1 - 10	243,546

In addition to the basic test files, two *bash* scripts were developed to test a large number of possible number ranges and to ensure that no deadlock was ever reached. The scripts run both solutions on a 100 by 100 matrix filled with values of all 1. Every possible M,N and K values for 1 to 100 is run for a total of 1 million test cases. Copies of these test scripts are included in Appendix A in addition to the source code. The multithreaded solution was run with *valgrind* to ensure that no memory leaks occurred during normal running or error situations.

Sample Input and Outputs

The section below indicates the input and output from the running of the several test files mentioned previously. For simplicity, only the last three lines of output has been shown.

Part A - Multiprocess

```
$ make
gcc -c pmms.c -Wall -Wextra -std=c99 -lrt -pthread -D _XOPEN_SOURCE=500
gcc -c fileIO.c -Wall -Wextra -std=c99 -lrt -pthread -D _XOPEN_SOURCE=500
gcc -c pmms.o fileIO.o -o pmms -Wall -Wextra -std=c99 -lrt -pthread -D
_XOPEN_SOURCE=500
```

```
$ ./pmms
Usage: ./pmms [MatrixA File] [MatrixB File] [M] [N] [K]
Please see README for detailed steps on how to run!
```

```
$ ./pmms allOnes.txt allOnes.txt 0 0 0
ERROR - matrix dimensions must be positive values
```

```
$ ./pmms allOnes.txt allOnes.txt 200 200 200
ERROR - not enough matrix values
```

```
$ ./pmms allOnes.txt allOnes.txt 100 100 100
Subtotal produced by process with ID 8120: 10000
Subtotal produced by process with ID 8121: 10000
Total: 1000000
```

```
$ ./pmms exampleA.txt exampleB.txt 3 2 4
Subtotal produced by process with ID 8623: 134
Subtotal produced by process with ID 8625: 206
Total: 402
```

```
$ ./pmms mediumA.txt mediumB.txt 10 10 10
Subtotal produced by process with ID 8694: 4671
Subtotal produced by process with ID 8689: 2829
Total: 27132
```

```
$ ./pmms largeA.txt largeB.txt 30 10 30
Subtotal produced by process with ID 8798: 8523
Subtotal produced by process with ID 8799: 8604
Total: 243546
```

Part B - Multithread

```
$ make
gcc -c pmms.c -Wall -Wextra -std=c99 -lrt -pthread -D _XOPEN_SOURCE=500
gcc -c fileIO.c -Wall -Wextra -std=c99 -lrt -pthread -D _XOPEN_SOURCE=500
gcc -c pmms.o fileIO.o -o pmms -Wall -Wextra -std=c99 -lrt -pthread -D
_XOPEN_SOURCE=500
```

```
$ ./pmms
Usage: ./pmms [MatrixA File] [MatrixB File] [M] [N] [K]
Please see README for detailed steps on how to run!
```

```
$ ./pmms allOnes.txt allOnes.txt 0 0 0
ERROR - matrix dimensions must be positive values
```

```
$ ./pmms allOnes.txt allOnes.txt 200 200 200
ERROR - not enough matrix values
```

```
$ ./pmms allOnes.txt allOnes.txt 100 100 100
Subtotal produced by thread with ID 139854480430848: 10000
Subtotal produced by thread with ID 139853857285888: 10000
Total: 1000000
```

```
$ ./pmms exampleA.txt exampleB.txt 3 2 4
Subtotal produced by thread with ID 140059911808768: 134
Subtotal produced by thread with ID 140059903416064: 206
Total: 402
```

```
$ ./pmms mediumA.txt mediumB.txt 10 10 10
Subtotal produced by thread with ID 140101363431168: 4671
Subtotal produced by thread with ID 140101463746304: 2829
Total: 27132
```

```
$ ./pmms largeA.txt largeB.txt 30 10 30
Subtotal produced by thread with ID 140287909222144: 8523
Subtotal produced by thread with ID 140287900829440: 8604
Total: 243546
```

References

Silberschatz, Abraham, Peter B. Galvin, and Greg Gagne. *Operating System Concepts*. 9th ed. Reading, MA: Addison-Wesley, 1994.

Soh, Sie Teng. "*Process and Threads*." Class lecture, Operating Systems from Curtin University, Perth, Australia, April 1 2016.

Soh, Sie Teng. "*Process Synchronization*." Class lecture, Operating Systems from Curtin University, Perth, Australia, April 15 2016.